

AI Infrastructure Options

VPS / Cloud Server · Cloud AI Platforms

Technical Specifications & Pricing Guide (INR)

Limitless Brain Lab | June 2026

OPTION 1

VPS / Linux Server

Rent a cloud VPS with GPU for local AI model inference — fixed monthly cost, full OS control.

OPTION 2

Cloud AI Platforms

Managed AI APIs (AWS Bedrock, Vertex AI, Anthropic) — zero server management, pay-per-token.

PART 1 — VPS / Linux Server (Ubuntu)

A VPS (Virtual Private Server) gives you a rented Linux server running 24/7 in a data centre. No upfront hardware purchase — you pay a fixed monthly fee. You get root/sudo access, install Ubuntu, and run your AI models exactly like a physical machine. Two VPS configurations are compared below.

Configuration A — GPU VPS (6 vCPU · 32 GB RAM · 16 GB VRAM · 500 GB SSD)

For running local AI models (LLaMA, Mistral, clinical LLMs) — requires dedicated GPU

Provider	Location	Specs	Monthly Cost (INR)	Notes
RunPod	Global / EU	RTX 4080 16GB VRAM 6 vCPU · 32GB RAM 500GB SSD	INR 28,000 – 36,000/mo	Best value GPU cloud. Spot = cheaper. Good for LLaMA 3 70B (int4)
Vast.ai	Global	RTX 4080 / 3090 16–24GB 6 vCPU · 32GB RAM 500GB SSD	INR 22,000 – 34,000/mo	Marketplace — lowest price but shared infra. Good for dev/test.
Lambda Labs	US / EU	A10 24GB VRAM 30 vCPU · 60GB RAM 1TB SSD	INR 52,000 – 68,000/mo	Reliable, dedicated GPU. Best for production inference workloads.
E2E Networks	India (Delhi/ Mumbai)	NVIDIA T4 16GB VRAM 8 vCPU · 32GB RAM 500GB SSD	INR 18,000 – 35,000/mo	Indian data centre — best for data residency in India. Govt. compliant.
AWS EC2 g4dn.2xlarge	ap-south-1 (Mumbai)	T4 16GB VRAM 8 vCPU · 32GB RAM 225GB NVMe	INR 40,000 – 58,000/mo	On-demand pricing. Reserve 1yr = ~40% cheaper (~INR 27,000/mo).
Google Cloud N1 + T4	asia-south1 (Mumbai)	T4 16GB VRAM 8 vCPU · 30GB RAM 500GB SSD	INR 38,000 – 54,000/mo	Preemptible / Spot VM = up to 60% cheaper (~INR 18,000/mo).
Azure NV6	Central India	NVIDIA M60 8GB VRAM 6 vCPU · 56GB RAM 380GB SSD	INR 35,000 – 48,000/mo	Good for inference. Spot pricing saves ~70% for dev workloads.

*Recommended for AI Inference: **RunPod** or **E2E Networks** — RunPod for best cost globally, E2E Networks if data must stay within India. For LLaMA 3 70B at int4 quantisation, 16GB VRAM is the minimum requirement.*

Configuration B — Standard VPS (4 vCPU · 16 GB RAM · Large Storage)

For web apps, databases, RAG backends, embedding servers — no GPU needed

Provider	Location	Specs	Monthly Cost (INR)	Notes
Hetzner CX32	Germany / Finland	4 vCPU · 8GB RAM 80GB SSD	INR 630 – 820/mo	Cheapest reliable VPS. Good for small apps and APIs.
Hetzner CPX41	Germany / Finland	8 vCPU · 16GB RAM 240GB SSD	INR 1,900 – 2,400/mo	Best price-to-performance. Ideal for RAG/embedding server.

DigitalOcean	Bangalore (India)	4 vCPU · 16GB RAM 320GB SSD	INR 6,500 – 8,500/mo	Indian region available. Simple dashboard, reliable uptime.
Linode (Akamai)	Mumbai (India)	4 vCPU · 16GB RAM 320GB SSD	INR 6,000 – 8,000/mo	Mumbai DC for India latency. Good support and SLA.
OVHcloud Advance-1	Singapore / EU	4-core Xeon · 32GB RAM 2×480GB SSD	INR 5,500 – 7,500/mo	Dedicated server (not shared vCPU). More reliable for DB workloads.
E2E Networks	India (Delhi/ Mumbai)	4 vCPU · 16GB RAM 200GB SSD	INR 4,000 – 6,000/mo	Indian data centre. Best for Indian compliance requirements.
Hostinger VPS	India / Singapore	4 vCPU · 16GB RAM 200GB NVMe	INR 1,500 – 2,500/mo	Budget-friendly. Good for small-medium projects. NVMe SSD.

*Recommended for Storage/Backend: **Hetzner CPX41 + Hetzner Storage Box** — unbeatable cost. For India data residency: **E2E Networks or DigitalOcean Bangalore**. For 200 TB storage, use object storage (Backblaze B2: ~INR 510/TB/mo or Wasabi: ~INR 600/TB/mo) instead of block storage VPS — block storage at 200 TB on any VPS would cost INR 5–15 lakh/month.*

Add-on: Storage Options for Large Data (INR / Month)

Storage Type	Provider	Capacity	Cost (INR/mo)	Best For
Object Storage	Backblaze B2 (India-accessible)	Any size	~INR 510/TB/mo	Backups, model weights, documents
Object Storage	Wasabi Hot Cloud (no egress fee)	Any size	~INR 600/TB/mo	Frequent access, media files, datasets
Object Storage	AWS S3 (ap-south-1 Mumbai)	Any size	~INR 2,000/TB/mo	Production apps needing AWS integration
Block Storage	Hetzner Volume (attached to VPS)	Up to 10 TB	~INR 450/100GB/mo	Database storage, pgvector index
NAS Storage Box	Hetzner Storage Box (SFTP/SMB)	Up to 20 TB	INR 1,000 – 2,800/mo	RAG document store, backup target
Object Storage	E2E Networks Object Storage (India)	Any size	~INR 400/TB/mo	India data residency, DPDP Act compliant

For 200 TB: Use Backblaze B2 or Wasabi — $200\text{ TB} \times \text{INR } 510 = \text{~INR } 1,02,000/\text{mo}$ (Backblaze B2). Far cheaper than block VPS storage. Mount via `s3fs` or `rclone` from your VPS.

PART 2 — Cloud AI Platforms

Cloud AI platforms let you call powerful AI models via API — no GPU server needed. You send text/data, get AI responses back. Pay only for what you use (per token). Three leading providers are evaluated below.

AWS Bedrock

Fully managed foundation model service on AWS

Key Models	Claude 3.5 Sonnet / Haiku, Llama 3.1 / 3.3, Mistral, Amazon Titan, Amazon Nova, Cohere Embed — all via single unified API.
Pricing (approx.)	Pay-per-token (no upfront cost): • Claude 3.5 Haiku: \$0.001 / 1K input · \$0.005 / 1K output • Claude 3.5 Sonnet: \$0.003 / 1K input · \$0.015 / 1K output • Llama 3.3 70B: \$0.00072 / 1K input · \$0.00099 / 1K output INR equivalent: approx. INR 500 – 1,200 per million tokens. Provisioned Throughput available for high volume at flat monthly rate.
UAE / MENA Region	Region: me-central-1 (UAE — Abu Dhabi). Data stays within UAE. Supports AWS PrivateLink. SOC 2, ISO 27001, PCI DSS certified.

Google Vertex AI

Unified ML platform with Gemini and 130+ partner models

Key Models	Gemini 2.0 Flash / Pro, Gemini 1.5 Flash / Pro, Imagen 3, Llama 3 (Model Garden), Mistral, Codey — plus AutoML and custom training.
Pricing (approx.)	Pay-per-token. Typical rates: • Gemini 1.5 Flash: \$0.000075 / 1K input · \$0.0003 / 1K output • Gemini 1.5 Pro: \$0.00125 / 1K input · \$0.005 / 1K output • Gemini 2.0 Flash: \$0.0001 / 1K input · \$0.0004 / 1K output INR equivalent: approx. INR 75 – 450 per million tokens. Committed Use Discounts (1yr/3yr) reduce cost by up to 57%.
UAE / MENA Region	Nearest region: me-central1 (Qatar). UAE region not yet GA. Supports VPC Service Controls. ISO 27001, SOC 2, PCI DSS, HIPAA eligible.

Anthropic Claude

Direct API + available via AWS Bedrock & Google Vertex AI

Key Models	Claude 4 Opus (highest intelligence), Claude 4 Sonnet (balanced), Claude 3.5 Haiku (fastest / cheapest) — 200K token context. Best for clinical document analysis and long-form reasoning.
Pricing (approx.)	Direct API (api.anthropic.com): • Claude 3.5 Haiku: \$0.001 / 1K input · \$0.005 / 1K output • Claude 3.5 Sonnet: \$0.003 / 1K input · \$0.015 / 1K output • Claude 4 Opus: \$0.015 / 1K input · \$0.075 / 1K output INR equivalent: approx. INR 500 – 6,200 per million tokens. Prompt caching = 90% cost reduction on repeated context.
UAE / MENA Region	No dedicated UAE DC on direct API. Use via AWS Bedrock (me-central-1 Abu Dhabi) for UAE data residency. GDPR + SOC 2 Type II. Data not used for training.

■ Your Data Is Safe & Secure — Complete Privacy Assurance

All three cloud AI providers — **AWS Bedrock, Google Vertex AI, and Anthropic Claude** — operate under enterprise-grade security. Your data, patient records, and business information are fully protected. **You do not need to worry about privacy or data security** when using any of these platforms.

■ Data Never Used for Training	AWS, Google, and Anthropic never use your API data to train or improve their models . Your inputs and outputs are completely private and isolated to your account only.
■ End-to-End Encryption	All data in transit: TLS 1.3 encryption . All data at rest: AES-256 encryption . Even cloud provider employees cannot read your data.
■ Data Residency Control	You choose which geographic region stores your data. AWS me-central-1 (Abu Dhabi) ensures data never leaves UAE borders.
■ Global Certifications	All providers hold SOC 2 Type II, ISO 27001, PCI DSS — the highest global standards for data security, audited annually by independent third parties.
■ Full Audit Trails	Every API call is logged with timestamps and user identity — full visibility for MOHAP, UAE PDPL, and DPDP Act compliance.
■ Complete Tenant Isolation	Your data is 100% isolated from other customers. No cross-tenant access possible. Role-based access controls (RBAC) ensure only authorised staff can call APIs.

Bottom line: All three cloud platforms are trusted by the world's largest hospitals, banks, and governments. Your data is in safe hands — protected by the same infrastructure that secures Fortune 500 companies globally.

UAE Data Privacy & Compliance

Regulation	Scope	Key Requirements	Cloud Impact
UAE PDPL (Fed. Law 45/2021)	All entities processing UAE residents' personal data	• Explicit consent for sensitive data • Data localisation for certain sectors • Right to access & erasure • Breach notification within 72 hrs	Use AWS me-central-1 (Abu Dhabi) for strict UAE data residency
DIFC DPL 2020	Companies in Dubai International Financial Centre	• GDPR-aligned framework • Data Protection Officer mandatory • Cross-border transfer restrictions • Privacy impact assessments required	AWS / GCP PrivateLink + VPC isolation recommended
ADGM DPR 2021	Entities in Abu Dhabi Global Market	• Based on UK GDPR • Lawful basis for processing • Data retention limits • Subject Access Requests within 30 days	Data must not leave UAE without adequacy decision
MOHAP Health Data	Healthcare providers & medical platforms in UAE	• Patient records classified as sensitive • Must be stored within UAE • Access limited to authorised clinical staff • Audit trails mandatory	AWS Bedrock (me-central-1) + AES-256 at rest + TLS 1.3 in transit

Recommendation: Use AWS Bedrock (me-central-1 / Abu Dhabi) for all cloud AI inference to guarantee UAE data residency. Enable PrivateLink so patient data never traverses the public internet.

Side-by-Side Comparison Summary

Factor	Config A GPU VPS	Config B Standard VPS	AWS Bedrock (Cloud AI)	Vertex AI (Cloud AI)	Anthropic (Cloud AI)
Data Control	Full — your VPS	Full — your VPS	High — UAE region	Medium — Qatar	Medium — no UAE DC
Setup Cost	Zero (no hardware)	Zero (no hardware)	Zero (PAYG)	Zero (PAYG)	Zero (PAYG)
Monthly Cost	INR 18,000–68,000/mo (GPU VPS)	INR 630–8,500/mo (no GPU)	Per-token variable	Per-token variable	Per-token variable
GPU / VRAM	16 GB VRAM (dedicated)	None (CPU only)	Managed by AWS	Managed by GCP	Managed by Anthropic
Model Choice	Any open-source model	Small / embed models	30+ models	130+ models	Claude family only
Max Context	Depends on VRAM	Limited	200K (Claude)	2M (Gemini)	200K tokens
UAE Compliant	Yes — choose India DC	Yes — choose India DC	Yes — me-central-1	Partial — Qatar	Via Bedrock/Vertex
Scalability	Fixed VPS capacity	Fixed VPS capacity	Instant auto-scale	Instant auto-scale	Instant auto-scale
Maintenance	You manage OS/software	You manage OS/software	Fully managed by AWS	Fully managed by GCP	Managed by Anthropic
Best For	Clinical AI, LLM chatbot	RAG backend, storage	Production AI API	Multi-modal, big context	Highest quality text AI

Recommended Architecture for Healthcare AI (UAE)

■ Primary AI Inference	Config A GPU VPS on E2E Networks (India) or RunPod — runs AI models, zero data leaves the server.
■ Overflow / Burst	AWS Bedrock (me-central-1) when VPS GPU is at capacity — clinical data stays in UAE region.
■ Storage / RAG	Config B VPS (Hetzner/E2E) + Backblaze B2 / Wasabi object storage for document store feeding AI.
■ Embeddings	Local embedding model on Config B VPS — avoids sending raw patient text to cloud providers.
■ Monitoring	Prometheus + Grafana on VPS; CloudWatch for Bedrock usage and cost alerts.
■ Compliance Audit	All inference requests logged to PostgreSQL with patient_id hash per MOHAP/UAE PDPL guidelines.

This document is a technical options guide prepared by Limitless Brain Lab (June 2026). VPS and cloud pricing is indicative and subject to change. All INR figures based on approx. USD/INR = 83.5. Regulatory guidance should be verified with qualified UAE/India legal counsel before deployment.